

# Prediction Project

Angelicaqj

5/10/2021

Predicting crime rates in Boston data. The Boston data set is in the MASS package, and first you will need to load it.



```
##      crim zn indus chas   nox    rm  age    dis rad tax ptratio  black lstat
## 1 0.00632 18  2.31    0 0.538 6.575 65.2 4.0900   1 296    15.3 396.90  4.98
## 2 0.02731  0  7.07    0 0.469 6.421 78.9 4.9671   2 242    17.8 396.90  9.14
## 3 0.02729  0  7.07    0 0.469 7.185 61.1 4.9671   2 242    17.8 392.83  4.03
## 4 0.03237  0  2.18    0 0.458 6.998 45.8 6.0622   3 222    18.7 394.63  2.94
## 5 0.06905  0  2.18    0 0.458 7.147 54.2 6.0622   3 222    18.7 396.90  5.33
## 6 0.02985  0  2.18    0 0.458 6.430 58.7 6.0622   3 222    18.7 394.12  5.21
##   medv
## 1 24.0
## 2 21.6
## 3 34.7
## 4 33.4
## 5 36.2
## 6 28.7
```

Build a regression model to predict the crime rate (crim) in Boston suburbs based on the other provided variables.

The solution includes:

- A brief exploratory analysis.
- A description of the set of regression models to be considered.
- A description of how the models were evaluated.
- A summary of one model, based on the analysis, is the best among to be considered.

## Exploratory Analysis



Generalize a linear model

```
##
## Call:
## glm(formula = crim ~ ., data = Boston)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -9.924  -2.120  -0.353   1.019  75.051
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  17.033228   7.234903   2.354 0.018949 *
## zn           0.044855   0.018734   2.394 0.017025 *
## indus        -0.063855   0.083407  -0.766 0.444294
## chas         -0.749134   1.180147  -0.635 0.525867
## nox          -10.313535   5.275536  -1.955 0.051152 .
## rm           0.430131   0.612830   0.702 0.483089
## age          0.001452   0.017925   0.081 0.935488
## dis          -0.987176   0.281817  -3.503 0.000502 ***
## rad           0.588209   0.088049   6.680 6.46e-11 ***
## tax          -0.003780   0.005156  -0.733 0.463793
## ptratio      -0.271081   0.186450  -1.454 0.146611
## black        -0.007538   0.003673  -2.052 0.040702 *
## lstat         0.126211   0.075725   1.667 0.096208 .
## medv         -0.198887   0.060516  -3.287 0.001087 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 41.46327)
##
##      Null deviance: 37363  on 505  degrees of freedom
## Residual deviance: 20400  on 492  degrees of freedom
## AIC: 3336.5
##
## Number of Fisher Scoring iterations: 2
```

The approach that we have been taking in this dataset is to use regression as a way of summarizing relationships between some of the variables in the dataset.



## Set Regression Models and evaluation

Best subset selection by identifying the best model that contain a given number of predictors.

```
## [1] 0.3912567 0.4207965 0.4286123 0.4334892 0.4392738 0.4440173 0.4476594
## [8] 0.4504606 0.4524408 0.4530572 0.4535605 0.4540031 0.4540104
```

It seems that the  $R^2$  statistic increases from 39%, when only one variable is included in the model, to 45%, when all variables are included. As expected, the  $R^2$  statistic increases monotonically as more variables are included.

```
## [1] 3
```



In the first plot, we see that there are three variables that share a BIC close to -260. These three variables are representing in the second plot as **rad**, **black**, and **lstat**, that contain the lowest BIC.

```
## Subset selection object
## Call: regsubsets.formula(crim ~ ., data = Boston, nvmax = 13, method = "forward")
## 13 Variables (and intercept)
##           Forced in Forced out
## zn          FALSE      FALSE
## indus        FALSE      FALSE
## chas          FALSE      FALSE
## nox           FALSE      FALSE
## rm            FALSE      FALSE
## age           FALSE      FALSE
## dis           FALSE      FALSE
## rad           FALSE      FALSE
## tax           FALSE      FALSE
## ptratio       FALSE      FALSE
## black         FALSE      FALSE
## lstat         FALSE      FALSE
## medv          FALSE      FALSE
## 1 subsets of each size up to 13
## Selection Algorithm: forward
##           zn  indus chas nox  rm  age dis  rad  tax ptratio black lstat medv
## 1  ( 1 )  " " " "  " " " " " " " " " " " " " " " " " " " " " " "
## 2  ( 1 )  " " " "  " " " " " " " " " " " " " " " " " " " " " "
## 3  ( 1 )  " " " "  " " " " " " " " " " " " " " " " " " " " " "
## 4  ( 1 )  " " " "  " " " " " " " " " " " " " " " " " " " " " "
## 5  ( 1 )  "*" " "  " " " " " " " " " " " " " " " " " " " " " "
## 6  ( 1 )  "*" " "  " " " " " " " " " " " " " " " " " " " " " "
## 7  ( 1 )  "*" " "  " " "*" " " " " " " " " " " " " " " " " " " "
## 8  ( 1 )  "*" " "  " " "*" " " " " " " " " " " " " " " " " " " "
## 9  ( 1 )  "*" "*"  " " "*" " " " " " " " " " " " " " " " " " " "
## 10 ( 1 )  "*" "*"  " " "*" "*" " " " " " " " " " " " " " " " " " "
## 11 ( 1 )  "*" "*"  " " "*" "*" " " " " " " " " " " " " " " " " " "
## 12 ( 1 )  "*" "*"  "*" "*" "*" " " " " " " " " " " " " " " " " " "
## 13 ( 1 )  "*" "*"  "*" "*" "*" "*" "*" "*" " " " " " " " " " " " "

## Subset selection object
## Call: regsubsets.formula(crim ~ ., data = Boston, nvmax = 13, method = "backward")
## 13 Variables (and intercept)
##           Forced in Forced out
## zn          FALSE      FALSE
## indus        FALSE      FALSE
## chas          FALSE      FALSE
## nox           FALSE      FALSE
## rm            FALSE      FALSE
## age           FALSE      FALSE
## dis           FALSE      FALSE
## rad           FALSE      FALSE
## tax           FALSE      FALSE
## ptratio       FALSE      FALSE
## black         FALSE      FALSE
## lstat         FALSE      FALSE
## medv          FALSE      FALSE
## 1 subsets of each size up to 13
## Selection Algorithm: backward
##           zn  indus chas nox  rm  age dis  rad  tax ptratio black lstat medv
```

We can see that using forward and backward stepwise selection, the best one-variable model contains only **rad**.

```
## (Intercept)          dis          rad          medv
##    3.4931998 -0.3241247  0.5152751 -0.1584437
```

```
## [1] 43.80667 43.80505
```

The best model according the Best Subset selection and Forward Stepwise selection is a model with three variables. The BIC plot shows noticeable that **rad**, **black**, and **lstat** have the lowest BIC. Concluding, the regression model to predict the crime rate (crim) in Boston suburbs is the following:

Based on the coefficients the final prediction model is the following:

```
##  
## Call:  
## lm(formula = crim ~ rad + black + lstat, data = Boston)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -11.023  -1.713  -0.281   0.873  77.716
```

```
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.372585   1.641557  -0.227  0.82054
## rad          0.488172   0.040422  12.077 < 2e-16 ***
## black       -0.009472   0.003615  -2.620  0.00905 **
## lstat        0.213596   0.047447   4.502 8.39e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.521 on 502 degrees of freedom
## Multiple R-squared:  0.4286, Adjusted R-squared:  0.4252
## F-statistic: 125.5 on 3 and 502 DF,  p-value: < 2.2e-16
```